


UNIVERSITY OF TWENTE.



NumPy

Array handling with NumPy



Presented by:
Mahdi Farnaghi
Assistant Professor,
Department of Geo-information Processing
ITC, The University of Twente





Why NumPy array

3

Array creation

7

Array Attributes

12

Array Indexing, Slicing, and Reshaping

14

Concatenation

25

Universal Functions

28

Aggregations

33

Broadcasting

36

Comparisons and Boolean Mask

39

Conclusion

45

Why NumPy array





Python and dynamic typing

- Python is well-known for its ease of use
- One of the characteristics: **DYNAMIC TYPING**
- In contrast to strongly typed languages, in Python we do not need to define the type of a variable
- But standard Python implementation is written in C
 - A C integer is a label for a memory position
 - A Python integer is pointer to an object with additional information
 - Additional information allows Python to be dynamic
 - This additional information result in additional cost, in particular for complex data structures

```
/* C code */  
int result = 0;  
for(int i=0; i<100; i++){  
    result += i;  
}
```

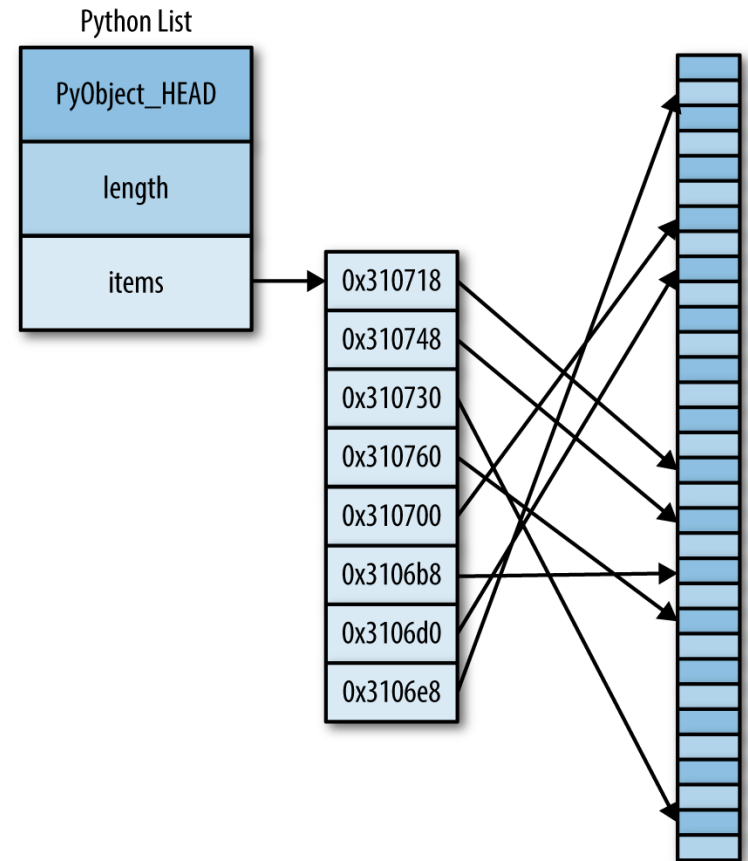
```
# Python code  
result = 0  
for i in range(100):  
    result += i
```





Python, dynamic typing, and Python List

- Python List
 - Each item is complete Python object
 - If items in the list are the same, then a lot of redundant data
 - Algorithms need to deal with the additional information of every object in the list
 - **Not Efficient**
- Python answer: Fixed-Type Arrays in Python
- A better solution: NumPy array, which is Fixed-Type and ...



Ref: Python data science handbook, by Jake VanderPlas

NumPy



Python library



Typed N-Dimensional
Arrays

Fast and versatile, vectorization,
indexing, and broadcasting
the de-facto standards of array
computing today.



Numerical computing

High-level mathematical functions,
random number generators, linear
algebra routines, Fourier
transforms, and more



Performant

Core: well-optimized C code;
Compiled; High speed



Interoperable

Supports a wide range of hardware
and computing platforms, and plays
well with distributed, GPU, and
sparse array libraries



Easy to use



Open source

Array creation



Numpy array creation

We assume throughout these slides that we have
import numpy as np

Array creation

- Conversion from another Python data structure: list or tuple

np.array([1, 2, 3]) or *np.array((6.05673, 52.56220))*

- Note: NumPy is constrained to arrays of the same type! NumPy will UPCAST if possible

array([3.14, 2, 3])

NumPy array creation

- Array generator functions

```
np.array([1,2,3])
```

```
np.arange(6) # half-open interval [0, 6)  
# numpy.arange([start, ]stop, [step, ]dtype=None)
```

```
np.linspace(0, 1, 12) # closed interval with number of steps  
# np.linspace(start, stop, num=50, endpoint=True, retstep=False, dtype=None)
```

```
np.zeros(10, dtype=float)  
np.ones((2, 3), dtype=float)  
np.full((4,2), 5) # Create a 4×2 array filled with 3.14
```

Other array generators: **np.random**, **np.normal**, **np.eye**, **np.empty**

Array Attributes

```
x1 = np.random.randint(5, size=6)
x2 = np.random.randint(10, size=(3,4))
x3 = np.random.randint(10, size=(2,3,4))
```

```
print("x1 ndim: ", x1.ndim)
print("x1 shape: ", x1.shape)
print("x1 size: ", x1.size)
print()
print("x2 ndim: ", x2.ndim)
print("x2 shape: ", x2.shape)
print("x2 size: ", x2.size)
print()
print("x3 ndim: ", x3.ndim)
print("x3 shape: ", x3.shape)
print("x3 size: ", x3.size)
```

```
x1 ndim: 1
x1 shape: (6,)
x1 size: 6
```

```
x2 ndim: 2
x2 shape: (3, 4)
x2 size: 12
```

```
x3 ndim: 3
x3 shape: (2, 3, 4)
x3 size: 24
```

```
x1
array([4, 0, 3, 3, 3, 1])
```

```
x3
array([[[8, 1, 5, 9],
        [8, 9, 4, 3],
        [0, 3, 5, 0]],
       [[2, 3, 8, 1],
        [3, 3, 3, 7],
        [0, 1, 9, 9]])
```

```
x2
array([[3, 5, 2, 4],
       [7, 6, 8, 8],
       [1, 6, 7, 7]])
```

NumPy array attributes

Four important attributes

- ndim
- shape
- size
- dtype

Array Indexing, Slicing, and Reshaping



Array indexing

- Quite similar with Python's standard list indexing
- Takes place between [and]
- Various expression types can be between the brackets
- Is allowed in left-hand and right-hand expression of an assignment

arrA[] = arrB[]

Single element indexing

- *arrA[4]* is the 5th element as it is 0-based
- *arrA[-2]* is the one but last element



Reshaping

```
arrA = np.arange(8)
```

```
arrA
```

```
array([0, 1, 2, 3, 4, 5, 6, 7])
```

```
arrA = arrA.reshape((2,4))
```

```
arrA
```

```
array([[0, 1, 2, 3],  
       [4, 5, 6, 7]])
```

```
arrA[1,3]
```

```
7
```

```
arrA[1]
```

```
array([4, 5, 6, 7])
```

just reshape it

no need to do [1][3]

arrA[1][3] still works but is inefficient

no index for each dimension, so

lower dimension array returned

Observe this ... in assignments

```
arrA[1,3] = 8
```

```
arrA
```

```
array([[0, 1, 2, 3],  
       [4, 5, 6, 8]])
```

```
arrA[1] = [4,5,6]
```

Traceback (most recent call last): File "<stdin>", line 1, in <module>

ValueError: cannot copy sequence with size 3 to array axis with dimension 4

```
arrA[:,1] = [10,12]
```

```
arrA
```

```
array([[ 0, 10,  2,  3],  
       [ 4, 12,  6,  8]])
```

works as long as dimension is right

Observe that `arrA` and
`arrA[:,1]` share memory
on values



Slicing

- In any dimension, a **slice** is a possibly empty list of indices
- Normal notation is $n : k$ standing for the list of indices from n to k , **n inclusive and k exclusive**.
- If we leave out n , it gets replaced by 0 (read it as "from start")
If we leave out k , it gets replaced by the highest index value possible (read it as "to end")
- Not often used but allowed is also **a stride value (aka step size)**, with notation $n : k : s$ standing for the list of indices going from n to k with step s .
- If we leave out s , we leave out $:$ also, and use 1 as step value.
- Slices and strides allow you to "cut through the data cube."
- Cf. *arange*'s three parameters.



Slicing

```
x = np.arange(10, 30)
```

```
x[::-1] # all elements, reversed
```

```
array([29, 28, 27, 26, 25, 24, 23, 22, 21, 20, 19, 18, 17, 16, 15, 14, 13,
       12, 11, 10])
```

```
x[::2] # every other element
```

```
array([10, 12, 14, 16, 18, 20, 22, 24, 26, 28])
```

- Multidimensional

```
x = np.arange(0, 24).reshape((4,6))
```

```
x
```

```
array([[ 0,  1,  2,  3,  4,  5],
       [ 6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17],
       [18, 19, 20, 21, 22, 23]])
```

```
x[:2,:3]
```

```
array([[0, 1, 2],
       [6, 7, 8]])
```



Slicing

`x[:,1]` # Second column

```
x[:,1] # Second column
```

`x[2,:]` # Third row

```
array([12, 13, 14, 15, 16, 17])
```

A subtle difference between NumPy array and Python list:
NumPy slicing returns **view** rather than *copy*

```
print(x)
x_slice = x[0, 4:]
x_slice[0] = 100
x_slice[1] = 200
print(x)
```

```
array([[ 0,  1,  2,  3,  4,  5],
       [ 6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17],
       [18, 19, 20, 21, 22, 23]])
```

```
array([[ 0,  1,  2,  3, 100, 200],
       [ 6,  7,  8,  9, 10, 11],
       [12, 13, 14, 15, 16, 17],
       [18, 19, 20, 21, 22, 23]])
```



Changing the shape of an array

- reshape function
- Add an axis with *np.newaxis* (this adds no values, just reshapes)
- ```
xx = np.arange(5)
xx[:, np.newaxis]
array([[0],
 [1],
 [2],
 [3],
 [4]])
```

# Concatenation

---



# Joining arrays

- Three important functions
  - `np.concatenate()`
  - `np.vstack()`
  - `np.hstack()`

```
x = np.array([4,5,6])
```

```
y = np.array([8,9,10])
```

```
np.concatenate([x, y])
```

```
array([4, 5, 6, 8, 9, 10])
```



# Joining arrays

- Concatenation along other axes

```
z = np.array([[11,22,33],
 [99,88,77]])
```

```
array([[11, 22, 33],
 [99, 88, 77]])
```

```
np.concatenate([z, z]) # concatenate along the first axis
```

```
array([[11, 22, 33],
 [99, 88, 77],
 [11, 22, 33],
 [99, 88, 77]])
```

```
np.concatenate([z, z], axis = 1) # concatenate along the second axis
```

```
array([[11, 22, 33, 11, 22, 33],
 [99, 88, 77, 99, 88, 77]])
```

# Universal Functions

---



# Universal Functions

- The importance of NumPy is related to its *vectorized operations*
- Using Python *List*, operations on arrays need to be done through *loop*  
e.g., imagine you want to compute the reciprocals of every element in a list, or you want to multiply every element by a constant value.
- Looping over Python List is very slow, due to the dynamic nature of the Python language
  - Type checking
  - Function dispatches
  - ...
- Solution
  - NumPy provides a convenient interface over several precompiled array operations
  - This precompiled array operations, known as ***vectorized*** operations, are accessible via NumPy ***ufuncs***
  - Extremely fast, particularly for big arrays
  - No need Loops





# UFuncs

- We apply Ufuncs to an array
- The function operates at element level
- With Python List, if we want to calculate the reciprocals of an array:

```
x = [2., 3., 5., 12., 30.]
reciprocals = []
for e in x:
 reciprocals.append(1. / e)
print(reciprocals)
```

- With Ufuncs:

```
x = np.array([2., 3., 5., 12., 30.], dtype=np.float)
reciprocals = 1. / x
print(reciprocals)
```

```
[0.5 0.33333333 0.2 0.08333333 0.03333333]
```



# UFuncs

```
x = np.arange(5)
```

We can directly apply Python arithmetic operators as UFuncs to NumPy arrays

```
x = np.arange(5)
print(.3 * x + x**2 + 12)
```

For each of these arithmetic operators there is a function counterpart implemented in numpy

| Operator | Name           |
|----------|----------------|
| +        | Addition       |
| -        | Subtraction    |
| *        | Multiplication |
| /        | Division       |
| %        | Modulus        |
| **       | Exponentiation |
| //       | Floor division |

```
array([0, 1, 2, 3, 4])
```

```
[12. 13.3 16.6 21.9 29.2]
```

|    |                 |
|----|-----------------|
| +  | np.add          |
| -  | np.subtract     |
| -  | np.negative     |
| *  | np.multiply     |
| /  | np.divide       |
| // | np.floor_divide |
| ** | np.power        |
| %  | np.mod          |



## Other UFuncs

- Trigonometric functions
- Hyperbolic functions
- Exponents and logarithms
- Rounding
- ...

```
sin(x, /[, out, where, casting, order, ...])
cos(x, /[, out, where, casting, order, ...])
tan(x, /[, out, where, casting, order, ...])
arcsin(x, /[, out, where, casting, order, ...])
arccos(x, /[, out, where, casting, order, ...])
arctan(x, /[, out, where, casting, order, ...])
```

```
around(a[, decimals, out])
```

```
round_(a[, decimals, out])
```

```
rint(x, /[, out, where, casting, order, ...])
```

```
fix(x[, out])
```

```
floor(x, /[, out, where, casting, order, ...])
```

```
ceil(x, /[, out, where, casting, order, ...])
```

```
trunc(x, /[, out, where, casting, order, ...])
```

```
exp(x, /[, out, where, casting, order, ...])
```

```
expm1(x, /[, out, where, casting, order, ...])
```

```
exp2(x, /[, out, where, casting, order, ...])
```

```
log(x, /[, out, where, casting, order, ...])
```

```
log10(x, /[, out, where, casting, order, ...])
```

```
log2(x, /[, out, where, casting, order, ...])
```

```
log1p(x, /[, out, where, casting, order, ...])
```

```
logaddexp(x1, x2, /[, out, where, casting, ...])
```

```
logaddexp2(x1, x2, /[, out, where, casting, ...])
```

- See <https://numpy.org/doc/stable/reference/routines.math.html> for a complete list of available functions

# Aggregations

---



# Aggregations

- Simple aggregation functions:  
`np.sum()`  
`np.min()`  
`np.max()`
- Multidimensional aggregations
  - E.g., how to calculate the maximum value of each row or column in a matrix  

```
x = np.arange(1, 16).reshape((5, 3))
print(x)
print("Sum of columns: ", x.sum(axis=0))
print("Sum of rows: ", x.sum(axis=1))
```

```
[[1 2 3]
 [4 5 6]
 [7 8 9]
 [10 11 12]
 [13 14 15]]
Sum of columns: [35 40 45]
Sum of rows: [6 15 24 33 42]
```

# List of NumPy aggregation functions

- There are also a NaN-safe version of each of these functions (except for `np.all` and `np.any`)
  - E.g., `np.nansum()`
- NaN: Not a number

| Functions                                           | Description                                                          |
|-----------------------------------------------------|----------------------------------------------------------------------|
| <code>np.mean()</code>                              | Compute the arithmetic mean along the specified axis.                |
| <code>np.std()</code>                               | Compute the standard deviation along the specified axis.             |
| <code>np.var()</code>                               | Compute the variance along the specified axis.                       |
| <code>np.sum()</code>                               | Sum of array elements over a given axis.                             |
| <code>np.prod()</code>                              | Return the product of array elements over a given axis.              |
| <code>np.cumsum()</code>                            | Return the cumulative sum of the elements along a given axis.        |
| <code>np.cumprod()</code>                           | Return the cumulative product of elements along a given axis.        |
| <code>np.min()</code> , <code>np.max()</code>       | Return the minimum / maximum of an array or minimum along an axis.   |
| <code>np.argmin()</code> , <code>np.argmax()</code> | Returns the indices of the minimum / maximum values along an axis    |
| <code>np.all()</code>                               | Test whether all array elements along a given axis evaluate to True. |
| <code>np.any()</code>                               | Test whether any array element along a given axis evaluates to True. |

# Broadcasting

---



# Broadcasting

- In NumPy, for arrays of the same size, binary operations perform on an element-by-element basis:

```
x = np.array([10, 20 , 30])
```

```
y = np.array([3, 2, 1])
```

```
print(x + y)
```

```
[13 22 31]
```

- What happens if we apply a binary operation on two array with different size?

```
a = np.arange(3)
```

```
print(a)
```

```
print()
```

```
b = np.arange(3)[: , np.newaxis]
```

```
print(b)
```

```
print()
```

```
print(a + b)
```

```
[0 1 2]
```

```
[[0]
```

```
[1]
```

```
[2]]
```

```
[[0 1 2]
```

```
[1 2 3]
```


```
[2 3 4]]
```





# Cascading

- NumPy tries to stretch and broadcast one of the arrays (or both of them) to match the shape of the other array. Then it applies the binary operation on the output arrays.



[https://www.tutorialspoint.com/numpy/numpy\\_broadcasting.htm](https://www.tutorialspoint.com/numpy/numpy_broadcasting.htm)

# Comparisons and Boolean Mask

---



# Comparison operators on NumPy arrays

- Applying a comparison operator on NumPy arrays result in a Boolean array
  - Boolean array: an array of Boolean values (True/False)
- Example

```
x = np.array([7, 8, 13, 80, 12, 1, 5])
print(x > 14)
```

```
[False False False True False False False]
```

- Comparison operators: ==, !=, <, <=, >, >=



## Boolean masks

Expression `myArr [ myBoolArr ]` uses another type of indexing with *rather different semantics*.

- The `myBoolArr` must have same shape or be broadcastable to `myArr`.
- Result is a one-dim array with `myArr` values found at positions where `myBoolArr` has value `True`.
- Common example: `myArr [ 2 < myArr & myArr < 10 ]`

Boolean masks can also be used for assignment purposes.



## Example

```
x = np.array([7, 8, 13, 80, 12, 1, 5])
b = x >= 13
print(b)
```



```
[False False True True False False False]
```

```
y = x[b]
print(y)
```



```
[13 80]
```

# Conclusion

---



## Some good references to read

- We only covered the most important aspects of NumPy
- For more information about NumPy functionalities see:
  - <https://numpy.org/devdocs/user/quickstart.html>
  - [https://www.w3schools.com/python/numpy\\_intro.asp](https://www.w3schools.com/python/numpy_intro.asp)
  - <https://www.tutorialspoint.com/numpy/index.htm>